

ggplot2: Going further in the tidyverse


Michael Friendly

Psych 6135

<https://friendly.github.io/6135/>

A larger view: Data science


- Data science treats statistics & data visualization as parts of a larger process
 - Data import: text files, data bases, web scraping, ...
 - Data cleaning → “tidy data”
 - Model building & visualization
 - Reproducible report writing



2

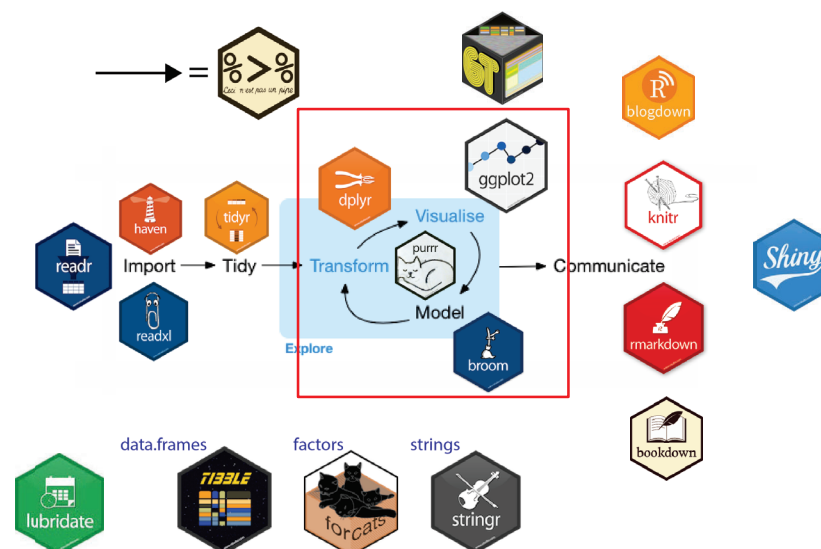


The tidyverse of R packages



3


The tidyverse expands



4

Topics

- Data import / export
- Data wrangling: getting your data into shape
 - dplyr & tidyr
 - pipes: %>%
 - grouping & summarizing
 - Example: NASA data on solar radiation
- Working with models: broom
 - Example: gapminder data
- Nice tables in R
- Bootstrapping



5

Ready for some heavy lifting?



6

Data Import / Export

- The readr package is the modern, tidy way to import and export data
 - Tabular data:
 - comma delimited (read.csv)
 - any other delimiters (";" = read.csv2; <tab> = read_tsv)
 - Data types:
 - specify column types or let functions guess
- Other data formats



package	Data types
haven	SAS, SPSS, Stata
readxl	Excel files (.xls and .xlsx)
DBI	Databases (SQL, ...)
rvest	HTML (web scraping)

7

Data Import: RStudio

Environment History Connections Build

Import Dataset

From Text (base)...
From Text (readr)...
From Excel...
From SPSS...
From SAS...
From Stata...

File/URL:
C:/Users/friendly/Dropbox/Documents/6135/R/drugs.txt

Data Preview:

subject (character)	drug1 (double)	drug2 (double)	drug3 (double)	drug4 (double)
sub1	20	34	38	44
sub2	16	28	30	34
sub3	14	28	26	30
sub4	18	20	24	30
sub5	10	18	14	22

options:

Import Options:

Name: drugs
Skip: 0

First Row as Names
Trim Spaces
Open Data Viewer

Delimiter: Whitespace
Quotes: Default
Locale: Configure...

Escape: None
Comment: Default
NA: Default

code:



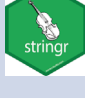
```
library(readr)
drugs <- read_table2("R/drugs.txt")
view(drugs)
```

Reading rectangular data using readr

8

Data transformation tools

Some common data types can be messy when imported. Tidy tools are there to help

dates/times	lubridate	read dates/times in various formats; extract components	
factors	forcats	Change order of levels, drop levels, combine levels	
strings	stringr	detect matches, subset, replace	

9



lubridate: Dates & times

PARSE DATE-TIMES (Convert strings or numbers to date-times)

1. Identify the order of the year (**y**), month (**m**), day (**d**), hour (**h**), minute (**m**) and second (**s**) elements in your data.
2. Use the function below whose name replicates the order. Each accepts a wide variety of input formats.

```

2017-11-28T14:02:00 ymd_hms(), ymd_hm(), ymd_h().
ymd_hms("2017-11-28T14:02:00")

2017-22-12 10:00:00 ydm_hms(), ydm_hm(), ydm_h().
ydm_hms("2017-22-12 10:00:00")

11/28/2017 1:02:03 mdy_hms(), mdy_hm(), mdy_h().
mdy_hms("11/28/2017 1:02:03")

1 Jan 2017 23:59:59 dmy_hms(), dmy_hm(), dmy_h().
dmy_hms("1 Jan 2017 23:59:59")

20170131 ymd(), ydm(). ymd(20170131)

July 4th, 2000 mdy(), myd(). mdy("July 4th, 2000")

4th of July '99 dmy(), dym(). dmy("4th of July '99")

2001: Q3 yq() Q for quarter. yq("2001: Q3")

2.01 hms:hms() Also lubridate:hms(),
hm() and ms(), which return
periods* hms: hms(sec = 0, min = 1,
hours = 2)
    
```

parse dates in various formats

```

ymd("20210604")
#> [1] "2021-06-04"
mdy("06-04-2021")
#> [1] "2021-06-04"
dmy("04/06/2021")
#> [1] "2021-06-04"
    
```

extract date components

```

minard_bday <- ymd("1781-03-27")
year(minard_bday)
#> [1] 1781
month.name[month(minard_bday)]
#> [1] "March"
    
```

date arithmetic: how old is Minard?

```

year(today()) - year(minard_bday)
#> [1] 241
    
```

Learn more at: <http://lubridate.tidyverse.org>

10



forcats: Working with factors

R represents categorical variables as factors, useful for analysis (e.g., ANOVA)
In graphics, we often want to recode levels or reorder them

Factors
R represents categorical data with factors. A factor is an integer vector with a levels attribute that stores a set of mappings between integers and categorical values. When you view a factor, R displays not the integers, but the values associated with them.

Create a factor with factor()
factor(x = character(), levels, labels = levels, exclude = NA, ordered = is.ordered(x), nmax = NA) Convert a vector to a factor. Also as_factor.
f <- factor(c("a", "c", "b", "a"), levels = c("a", "b", "c"))

Return its levels with levels()
levels(x) Return/set the levels of a factor. levels(f); levels(f) <- c("x", "y", "z")
Use unclass() to see its structure

Inspect Factors

fact_count(f, sort = FALSE)
Count the number of values with each level. fact_count(f)

Change the order of levels

fct_relevel(f, ..., after = 0L)
Manually reorder factor levels.
fct_relevel(c("b", "c", "a"))

fct_infreq(f, ordered = NA)
Reorder levels by the frequency in which they appear in the data (highest frequency first).
f <- factor(c("c", "c", "c", "a"))
fct_infreq(f)

fct_inorder(f, ordered = NA)
Reorder levels by order in which they appear in the data.
fct_inorder(f)

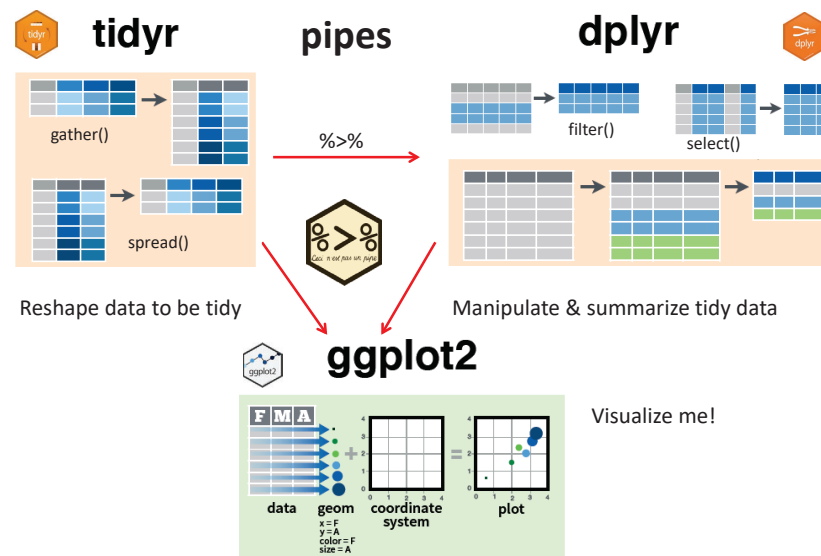
fct_rev(f) Reverse level order.
f <- factor(c("a", "b", "c"))
fct_rev(f)

fct_shift(f) Shift levels to left or right, wrapping around end.
fct_shift(f)

Learn more at: <http://forcats.tidyverse.org>

12

Tidy tools: overview



13

Tidy operations

Reshape wide to long
synonym: `tidyr::pivot_longer()`

Reshape long to wide
synonym: `tidyr::pivot_wider()`

`tidyr::gather(cases, "year", "n", 2:4)`
Gather columns into rows.

`tidyr::spread(pollution, size, amount)`
Spread rows into columns.

`tidyr::separate(storms, date, c("y", "m", "d"))`
Separate one column into several.

`tidyr::unite(data, col, ..., sep)`
Unite several columns into one.

Separate parts of a value into several variables

Join related variables into one


14

Data wrangling with dplyr & tidyr

What is Tidy Data?

A dataset is said to be tidy if:

- observations are in **rows**
- variables are in **columns**
- each value is in its own **cell**.



A “messy” dataset: Survey of income by religion from Pew Research

- Values of **income** are in separate columns, not one variable
- Column headers are **values**, not variable names
- Cell values are frequencies--- **implicit**, not explicit

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	\$50-75k
Agnostic	27	34	60	81	76	137
Atheist	12	27	37	52	35	70
Buddhist	27	21	30	34	33	58
Catholic	418	617	732	670	638	1116

This organization is easy in Excel

But, this makes data analysis and graphing hard

15

Tidying: reshaping wide to long

We can tidy the data by reshaping from wide to long format using `tidyr::gather()`

```
> pew <- read.delim(
  file = "http://stat405.had.co.nz/data/pew.txt",
  header = TRUE,
  stringsAsFactors = FALSE, check.names = FALSE)

> (pew1 <- pew[1:4, 1:6]) # small subset
```

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k
1 Agnostic	27	34	60	81	76
2 Atheist	12	27	37	52	35
3 Buddhist	27	21	30	34	33
4 Catholic	418	617	732	670	638

```
> library(tidyr)
> gather(pew1, "income", "frequency", 2:6)
```

	key	value	columns
1	religion	income	frequency
2	1 Agnostic	<\$10k	27
3	2 Atheist	<\$10k	12
4	3 Buddhist	<\$10k	27
5	4 Catholic	<\$10k	418
6	1 Agnostic	\$10-20k	34
7	2 Atheist	\$10-20k	27
8	3 Buddhist	\$10-20k	21
9	4 Catholic	\$10-20k	617
10	1 Agnostic	\$20-30k	60
11	2 Atheist	\$20-30k	37
12	3 Buddhist	\$20-30k	30
13	4 Catholic	\$20-30k	732
14	1 Agnostic	\$30-40k	81
15	2 Atheist	\$30-40k	52
16	3 Buddhist	\$30-40k	34
17	4 Catholic	\$30-40k	670

Another solution, using `reshape2::melt()`

```
> library(reshape2)
> pew_tidy <- melt(
  data = pew1,
  id = "religion",
  variable.name = "income",
  value.name = "frequency"
)
```

NB: income is a **character** variable; we might want to create an **ordered factor** or **numeric** version

16



Using pipes: %>%

• R is a functional language

- This means that $f(x)$ returns a value, as in $y <- f(x)$
- That value can be passed to another function: $g(f(x))$
- And so on: $h(g(f(x)))$

```
> x <- c(0.109, 0.359, 0.63, 0.996, 0.515, 0.142)
> exp(diff(log(x)))
[1] 3.29 1.75 1.58 0.52 0.28
```

- This gets messy and hard to read, unless you break it down step by step

```
> # Compute the logarithm of 'x', calculate lagged differences,
> # return the exponential function of the result
> log(x)
[1] -2.216 -1.024 -0.462 -0.004 -0.664 -1.952
> diff(log(x)) #calculate lagged diffs
[1] 1.19 0.56 0.46 -0.66 -1.29
> exp(diff(log(x))) # convert back to original scale
[1] 3.29 1.75 1.58 0.52 0.28
```

17

Using pipes: %>%

- Pipes (%>%) change the syntax to make this easier

```
> # use pipes
> x %>% log() %>% diff() %>% exp()
[1] 3.29 1.75 1.58 0.52 0.28
```

- Basic rules

- x %>% f() passes object on left hand side as first argument (or argument) of function on right hand side
 - x %>% f() is the same as f(x)
 - x %>% f(y) is the same as f(x, y)
 - y %>% f(x, ., z) is the same as f(x, y, z)
 - x %<% f() does the same, but assigns the result to x
 - Shortcut for x <- x %>% f()

18



Using pipes: %>% ggplot()




For the Pew data, mutate income → ordered factor and make a ggplot

```
pew1 %>%
  gather("income", "frequency", 2:6) %>% # reshape
  mutate(income = ordered(income, levels=unique(income))) %>% # make ordered
  ggplot(aes(x=income, fill=religion)) + # plot
  geom_bar(aes(weight=frequency)) # as freq bars
```

mutate() calculates or transforms column variables
ordered(income) levels are now ordered appropriately.

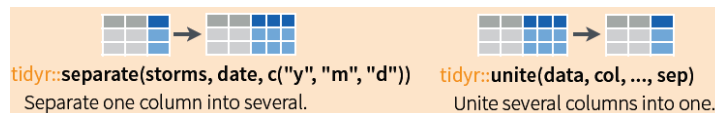
The result is piped to ggplot()



19

Tidying: separate() and unite()

It sometimes happens that several variables are crammed into one column, or parts of one variable are split across multiple columns



For example, for the pew data, we might want separate income into low & high

```
pew_long %>%
  mutate(inc = gsub("[\\$k]", "", income)) %>%
  mutate(inc = gsub("<", "0-", inc)) %>%
  separate(inc, c("low", "high"), ".") %>%
  head()
```

	religion	income	frequency	low	high
1	Agnostic	<\$10k	27	0	10
2	Atheist	<\$10k	12	0	10
3	Buddhist	<\$10k	27	0	10
4	Catholic	<\$10k	418	0	10
5	Agnostic	\$10-20k	34	10	20
6	Atheist	\$10-20k	27	10	20

20

dplyr: Subset observations (rows)

dplyr implements a variety of verbs to select a subset of observations from a dataset



dplyr::filter(iris, Sepal.Length > 7)
Extract rows that meet logical criteria.

dplyr::distinct(iris)
Remove duplicate rows.

dplyr::sample_frac(iris, 0.5, replace = TRUE)
Randomly select fraction of rows.

dplyr::sample_n(iris, 10, replace = TRUE)
Randomly select n rows.

dplyr::slice(iris, 10:15)
Select rows by position.

dplyr::top_n(storms, 2, date)
Select and order top n entries (by group if grouped data).

In a pipe expression, omit the dataset name

```
iris %>% filter(Sepal.Length > 7)
iris %>% filter(Species=="setosa")
```

```
iris %>% sample_n(10)
iris %>% slice(1:50) # setosa
```

21

dplyr: Subset variables (columns)



dplyr::select(iris, Sepal.Width, Petal.Length, Species)

Select columns by name or helper function.

Many helper functions in dplyr allow selection by a **function** of variable names:

select(iris, contains(""))

Select columns whose name contains a character string.

select(iris, ends_with("Length"))

Select columns whose name ends with a character string.

select(iris, everything())

Select every column.

select(iris, matches("t."))

Select columns whose name matches a regular expression.

select(iris, num_range("x", 1:5))

Select columns named x1, x2, x3, x4, x5.

select(iris, one_of(c("Species", "Genus")))

Select columns whose names are in a group of names.

select(iris, starts_with("Sepal"))

Select columns whose name starts with a character string.

select(iris, Sepal.Length:Petal.Width)

Select all columns between Sepal.Length and Petal.Width (inclusive).

select(iris, -Species)

Select all columns except Species.


22

dplyr: group_by() and summarise()

- Fundamental operations in data munging are:

- grouping** a dataset by one or more variables
- calculating one or more **summary** measures
- ungrouping**: expand to an ungrouped copy, if needed

data group summarise



mtcars %>%

group_by(cyl) %>%

summarise(avg=mean(mpg))

data avg



mtcars %>%

group_by(cyl) %>%

summarise(avg=mean(mpg)) %>%
ungroup()

23

Example: NASA data on solar radiation



Surface meteorology and Solar Energy

A renewable energy resource web site (release 6.0)
sponsored by NASA's Applied Science Program in the Science Mission Directorate
developed by POWER: Prediction of Worldwide Energy Resource Project




- over 200 satellite-derived meteorology and solar energy parameters
- monthly averaged from 22 years of data
- data tables for a particular location
- GIS Web Mapping Application & Services

How does solar radiation vary with latitude, over months of the year?

How to make this plot?

Q:
what are the basic plot elements?




24

NASA data: solar radiation

This is easy to do for the total **Annual** solar radiation, a column in the data

```
> str(nasa)
'data.frame': 64800 obs. of 15 variables:
 $ Lat: int -90 -90 -90 -90 -90 -90 -90 -90 -90 ...
 $ Lon: int -180 -179 -178 -177 -176 -175 -174 -173 -172 ...
 $ Jan: num 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 ...
 $ Feb: num 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 ...
 $ Mar: num 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 ...
 $ Apr: num 0 0 0 0 0 0 0 0 0 ...
 $ May: num 0 0 0 0 0 0 0 0 0 ...
 $ Jun: num 0 0 0 0 0 0 0 0 0 ...
 $ Jul: num 0 0 0 0 0 0 0 0 0 ...
 $ Aug: num 0 0 0 0 0 0 0 0 0 ...
 $ Sep: num 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 ...
 $ Oct: num 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 ...
 $ Nov: num 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 ...
 $ Dec: num 11 11 11 11 11 11 11 11 11 ...
 $ Ann: num 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 ...
```



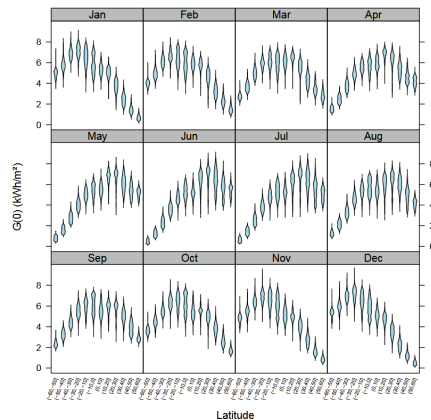
```
nasa %>%
  filter(abs(Lat) < 60) %>%
  mutate(Latf = cut(Lat, pretty(Lat, n=10))) %>%
  ggplot(aes(x=Latf, y=Ann)) +
  geom_violin(fill="pink", alpha=0.3) +
  labs(x="Latitude", y="Solar radiation G(0) (kWh/m²)")
```

25

Faceting & tidy data

This is complicated to do for the separate months, because the data structure is **untidy**--- months were in separate variables (wide format)

```
> str(nasa)
'data.frame': 64800 obs. of 15 variables:
 $ Lat: int -90 -90 -90 -90 -90 -90 -90 -90 ...
 $ Lon: int -180 -179 -178 -177 -176 -175 -174 -173 -172 -171 ...
 $ Jan: num 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 ...
 $ Feb: num 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 ...
 $ Mar: num 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 ...
 $ Apr: num 0 0 0 0 0 0 0 0 0 0 ...
 $ May: num 0 0 0 0 0 0 0 0 0 0 ...
 $ Jun: num 0 0 0 0 0 0 0 0 0 0 ...
 $ Jul: num 0 0 0 0 0 0 0 0 0 0 ...
 $ Aug: num 0 0 0 0 0 0 0 0 0 0 ...
 $ Sep: num 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 ...
 $ Oct: num 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 ...
 $ Nov: num 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 ...
 $ Dec: num 11 11 11 11 11 11 11 11 11 11 ...
 $ Ann: num 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 ...
```



26

tidying the data

To plot solar radiation against latitude by month (separate panels), we need to:

- remove the Ann column
- reshape the data to long format, so solar is all in one column

```
library(tidyr)
library(dplyr)
library(ggplot2)
```

```
nasa_long <- nasa %>%
  select(-Ann) %>%
  gather(month, solar, Jan:Dec, factor_key=TRUE) %>%
  filter(abs(Lat) < 60) %>%
  mutate(Lat_f = cut(Lat, pretty(Lat, 12)))
```

%>% “pipes” data to the next stage

select() extracts or drops columns

gather() collapses columns into key-value pairs

filter() subsets observations

mutate() creates new variables

27

tidying the data

```
> str(nasa_long)
'data.frame': 514080 obs. of 5 variables:
 $ Lat : int -59 -59 -59 -59 -59 -59 -59 -59 -59 ...
 $ Lon : int -180 -179 -178 -177 -176 -175 -174 -173 -172 -171 ...
 $ month: Factor w/ 12 levels "Jan","Feb","Mar",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ solar: num 5.19 5.19 5.25 5.25 5.17 5.17 5.15 5.15 5.15 5.15 ...
 $ Lat_f: Factor w/ 12 levels "(-60,-50]",...: 1 1 1 1 1 1 1 1 1 1 ...
```

For ease of plotting, I created a factor version of Lat with 12 levels

```
> head(nasa_long)
  Lat Lon month solar Lat_f
1 -59 -180 Jan 5.19 (-60,-50]
2 -59 -179 Jan 5.19 (-60,-50]
3 -59 -178 Jan 5.25 (-60,-50]
4 -59 -177 Jan 5.25 (-60,-50]
5 -59 -176 Jan 5.17 (-60,-50]
6 -59 -175 Jan 5.17 (-60,-50]
```

The data are now in a form where I can plot solar against Lat or Lat_f and facet by month

28

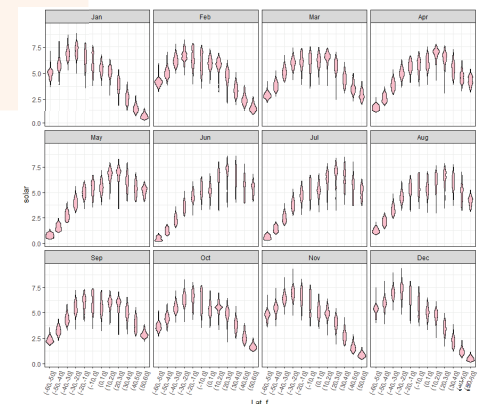
plotting the tidy data

Using **geom_violin()** shows the shapes of the distributions for levels of Lat_f

```
ggplot(nasa_long, aes(x=Lat_f, y=solar)) +
  geom_violin(fill="pink") +
  facet_wrap(~ month) +
  theme_bw() +
  theme(axis.text.x =
    element_text(angle = 70,
      hjust = 1))
```

facet_wrap(~month) does the right thing

I had to adjust the x-axis labels for Lat_f to avoid overplotting



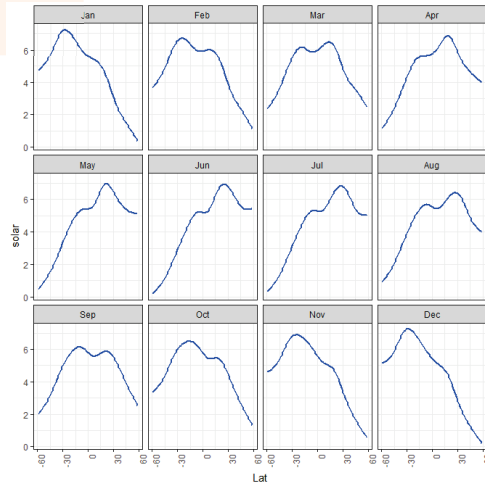
plotting the tidy data: smoothing

```
ggplot(nasa_long, aes(x=Lat, y=solar)) +
  geom_smooth(color="blue") +
  facet_wrap(~ month) +
  theme_bw()
```

Here we treat Lat as quantitative.
geom_smooth() uses method =
"gam" here because of large n

The variation in the smoothed
trends over the year suggest
quite lawful behavior

Can we express this as a statistical
model ?



build a model

What we saw in the plot suggests a **generalized additive model**, with a smooth, $s(\text{Lat})$

```
library(mgcv)
nasa.gam <- gam(solar ~ Lon + month + s(Lat), data=nasa_long)
summary(nasa.gam)
```

Family: gaussian
Link function: identity

Formula:
solar ~ Lon + month + s(Lat)

Parametric coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 4.691e+00 6.833e-03 686.409 < 2e-16 ***
Lon -1.713e-04 1.898e-05 -9.022 < 2e-16 ***
monthFeb 1.195e-01 9.664e-03 12.364 < 2e-16 ***
...
monthDec -8.046e-02 9.664e-03 -8.326 < 2e-16 ***
...
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The violin plots suggest that variance is not
constant. I'm ignoring this here by using the
default gaussian model.

Model terms:

- Lon wasn't included before
- month is a factor, for the plots
- $s(\text{Lat})$ fits a **smoothed term** in latitude,
averaged over other factors

There are other model choices, but it is
useful to visualize what we have done so
far


31

visualize the model


Effect plots show the fitted relationship between the response and model terms,
averaged over other predictors.

The {mgcv} package has its own versions of these.


```
plot(nasa.gam, cex.lab=1.25)
termplot(nasa.gam, terms="month", se=TRUE, lwd.term=3, lwd.se=2, cex.lab=1.25)
termplot(nasa.gam, terms="Lon", se=TRUE, lwd.term=3, lwd.se=2, cex.lab=1.25)
```



why the dip at the equator?



effect of longitude is very
small, but maybe
interpretable



month should be modeled
as a cyclic time variable

32

Visualizing models

- R modeling functions [lm(), glm(), ...] return model objects,
but these are "messy"
 - extracting coefficients takes several steps: data.frame(coef(mymod))
 - some info (R^2 , F , p .value) is computed in print() method, not stored
 - can't easily combine models
- Some have associated plotting functions
 - plot(model): diagnostic plots
 - car package: many model plot methods
 - effects package: plot effects for model terms
- But what if you want to:
 - make a table of model summary statistics
 - fit a **collection** of models, compare, summarize or visualize them?

33



broom: manipulating models

- The broom package turns model objects into tidy data frames
 - glance(models)** extracts model-level summary statistics (R^2 , df, AIC, BIC)
 - tidy(models)** extracts coefficients, SE, p-values
 - augment(models)** extracts observation-level info (residuals, ...)

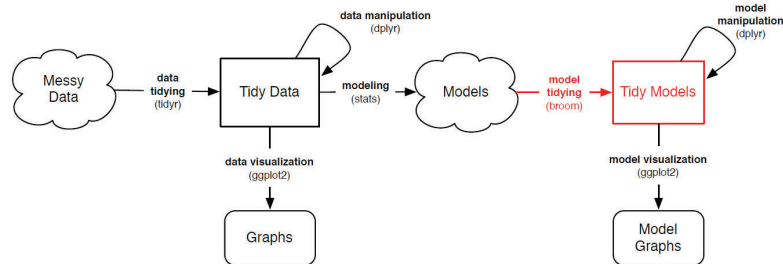


Image from: https://opr.princeton.edu/workshops/Downloads/2016Jan_BroomRobinson.pdf

34


Example: gapminder data

```
ggplot(aes(x = log(gdpPercap), y=lifeExp, color=continent), data=gapminder) +
  geom_point() +
  geom_smooth(method = "loess")
```

How to model this?

How to extract & plot model statistics?

How to fit & display multiple models for subsets?



35

Example: gapminder data

Predict life expectancy from year, population, GDP and continent:

```
gapmod <- lm(lifeExp ~ year + pop + log(gdpPercap) + continent, data=gapminder)
summary(gapmod)
```

```
Call:
lm(formula = lifeExp ~ year + pop + log(gdpPercap) + continent, data = gapminder)

Residuals:
    Min       1Q   Median       3Q      Max
-24.928  -3.285   0.314   3.699  15.221

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  -4.58e+02  1.67e+01  -27.43 < 2e-16 ***
year          2.38e-01  8.61e-03   27.58 < 2e-16 ***
pop           5.40e-09  1.38e-09    3.91 9.5e-05 ***
log(gdpPercap) 5.10e+00  1.60e-01   31.88 < 2e-16 ***
continentAmericas 8.74e+00  4.63e-01   18.86 < 2e-16 ***
continentAsia    6.64e+00  4.09e-01   16.22 < 2e-16 ***
continentEurope  1.23e+01  5.10e-01   24.11 < 2e-16 ***
continentOceania 1.26e+01  1.27e+00    9.88 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.79 on 1696 degrees of freedom
Multiple R-squared:  0.8,    Adjusted R-squared:  0.799
F-statistic: 969 on 7 and 1696 DF,  p-value: <2e-16
```

observation level

component level
(coefficients)

model level

36

glance() gives the **model level** summary statistics

```
> glance(gapmod)
  r.squared adj.r.squared sigma statistic p.value df logLik   AIC    BIC deviance df.residual
1      0.8      0.7992 5.789      969      0 8   -5406 10830 10879   56835      1696
```

tidy() gives the **model component** (term) statistics

```
> tidy(gapmod)
  term      estimate std.error statistic    p.value
1 (Intercept) -4.585e+02 1.671e+01  -27.433 1.982e-137
2 year         2.376e-01 8.613e-03   27.584 1.122e-138
3 pop          5.403e-09 1.381e-09    3.912 9.496e-05
4 log(gdpPercap) 5.103e+00 1.601e-01   31.876 4.096e-175
5 continentAmericas 8.739e+00 4.635e-01   18.856 3.758e-72
6 continentAsia    6.635e+00 4.091e-01   16.219 4.167e-55
7 continentEurope  1.230e+01 5.102e-01   24.113 1.044e-110
8 continentOceania 1.256e+01 1.270e+00    9.884 1.943e-22
```


augment() gives the **observation level** statistics

```
> augment(gapmod) %>% slice(1:5)
# A tibble: 5 x 12
  lifeExp year  pop log.gdpPercap. continent .fitted .se.fit .resid .hat .sigma
  <dbl> <int> <int> <dbl> <fct> <dbl> <dbl> <dbl> <dbl> <dbl>
1  28.8  1952  8425333 6.66 Asia 46.0 0.408 -17.1 0.00496 5.78
2  30.3  1957  9240934 6.71 Asia 47.4 0.390 -17.1 0.00454 5.78
3  32.0  1962 10267083 6.75 Asia 48.8 0.376 -16.8 0.00423 5.78
4  34.0  1967 11537966 6.73 Asia 49.9 0.372 -15.9 0.00413 5.78
5  36.1  1972 13079460 6.61 Asia 50.5 0.382 -14.4 0.00435 5.78
# ... with 2 more variables: .cooks <dbl>, .std.resid <dbl>
```

37

tidyr:: “nest – map – unnest” trick

- In many cases, we want to perform analysis for each subset of a dataset defined by one or more variables
 - dplyr::group_by(), summarise(), ungroup() is one way
- tidyr::nest(), purrr::map(), tidyr::unnest() is more general

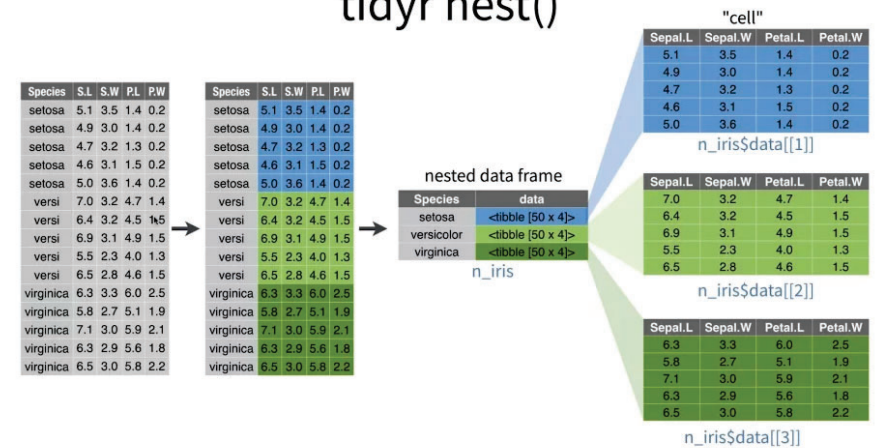


See: https://cran.r-project.org/web/packages/broom/vignettes/broom_and_dplyr.html

38

```
n_iris <- iris %>% group_by(Species) %>% nest() # group by Species, then nest
n_iris <- iris %>% nest(-Species)             # nest all other cols
```

tidyr nest()



CC BY RStudio

39


purrr::map() & friends



A fundamental operation is doing something, `f()`, to each element of a vector, list, or column of a data frame

`map(1:3, log) <=> list(log(1), log(2), log(3))`

`map(x, f)` returns a **list** of `f()` applied to each of `x`
Other variants, `map_dbl`, `int`, `chr` return **vectors**



data_frame : df

a	b
1	10
2	20
3	30
4	40

list : return

mean()
mean()

`map(df, mean)`

40

tidyr: fitting multiple models

There may be different effects by continent (GDP x continent interaction)

- What if want to fit (and visualize) a separate model for each continent?
- nest by continent, then {fit, tidy, glance, augment}

```
models <- gapminder %>%
  filter(continent != "Oceania") %>% # only two countries
  nest(data = -continent) %>%
  mutate(
    fit = map(data, ~ lm(lifeExp ~ year + pop + log(gdpPercap), data = .x)),
    tidied = map(fit, tidy),
    glanced = map(fit, glance),
    augmented = map(fit, augment)
  )
```

What's in this object?

```
names(models)
```

```
[1] "continent" "data"      "fit"       "tidied"    "glanced"   "augmented"
```

41

```
# view model summaries
```

```
models %>%
  select(continent, glanced) %>%
  unnest(glanced)
```

Model summary
statistics

```
# A tibble: 4 x 13
  continent r.squared adj.r.squared sigma statistic p.value df logLik AIC
  <fct>      <dbl>      <dbl> <dbl>      <dbl> <dbl> <dbl> <dbl> <dbl>
1 Asia      0.696      0.694 6.56      299. 5.27e-101 3 -1305. 2620.
2 Europe    0.797      0.795 2.46      466. 7.42e-123 3 -833. 1675.
3 Africa    0.500      0.498 6.48      207. 5.90e-93 3 -2050. 4110.
4 Americas  0.720      0.718 4.97      254. 1.39e-81 3 -904. 1819.
# ... with 4 more variables: BIC <dbl>, deviance <dbl>, df.residual <int>,
```

```
# model coefficients & tests
```

```
models %>%
  select(continent, tidied) %>%
  unnest(tidied)
```

Coefficients

```
# A tibble: 16 x 6
  continent term      estimate std.error statistic p.value
  <fct>      <chr>      <dbl>      <dbl>      <dbl>      <dbl>
1 Asia      (Intercept) -6.20e+2  4.00e+1  -15.5    1.34e-42
2 Asia      year        3.23e-1  2.06e-2   15.7    2.41e-43
3 Asia      pop         5.13e-9  1.66e-9    3.09    2.15e-3
4 Asia      log(gdpPercap) 5.04e+0  2.76e-1   18.3    2.25e-54
5 Europe    (Intercept) -1.72e+2  1.72e+1  -10.0    4.51e-21
# ...
```

42

```
# observation-level statistics
```

```
models %>%
  select(continent, augmented) %>%
  unnest(augmented)
```

Observations

```
# A tibble: 1,680 x 10
  continent lifeExp year      pop `log(gdpPercap)` .fitted .hat .sigma .cooksd
  <fct>      <dbl> <int>      <int>      <dbl>      <dbl> <dbl> <dbl> <dbl>
1 Asia      28.8  1952  8425333      6.66      43.7 0.0101 6.53 0.0133
2 Asia      30.3  1957  9240934      6.71      45.6 0.00822 6.53 0.0113
3 Asia      32.0  1962  10267083      6.75      47.4 0.00685 6.53 0.00957
4 Asia      34.0  1967  11537966      6.73      48.9 0.00616 6.53 0.00805
5 Asia      36.1  1972  13079460      6.61      49.9 0.00645 6.54 0.00727
6 Asia      38.4  1977  14880372      6.67      51.9 0.00640 6.54 0.00678
7 Asia      39.9  1982  12881816      6.89      54.6 0.00607 6.53 0.00771
8 Asia      40.8  1987  13867957      6.75      55.5 0.00795 6.53 0.0101
9 Asia      41.7  1992  16317921      6.48      55.8 0.0114 6.53 0.0134
10 Asia     41.8  1997  22227415      6.45      57.3 0.0138 6.53 0.0198
# ... with 1,670 more rows, and 1 more variable: .std.resid <dbl>
```


y predictors y diagnostics

43

Visualizing multiple models

One visual summary might be a plot of R^2 values, ordered by continent

```
models %>%
  select(continent, glanced) %>% unnest(glanced) %>%
  ggplot(aes(r.squared, reorder(continent, r.squared))) +
  geom_point(size=4) +
  geom_segment(aes(xend = 0, yend = ..y..)) +
  ylab("Continent")
```



44

Visualizing coefficients


Coefficient plots are often useful, but these are on different scales.

```
models %>% select(continent, tidied) %>% unnest(tidied)
filter(term != "(Intercept)") %>%
mutate(term=factor(term, levels=c("log(gdpPercap)", "year", "pop"))) %>%
ggplot(aes(x=term, y=statistic, color=continent, group=continent)) +
  geom_point(size=5, alpha=0.5) +
  geom_line(size=1.5) +
  geom_hline(yintercept=c(-2, 0, 2), color = c("red", "black", "red")) +
  ylab("t statistic") +
  theme_minimal() + theme(legend.position=c(0.9, 0.8))
```

get model stats
ignore the intercept
reorder terms sensibly
hlines for non-significance

Here, I plot the t -statistics, $t = b_{ij}/se(b_{ij})$ for all terms in all models.

Any values outside $\sim \pm 2$ are significant, $p < 0.5$!




45

Visualizing model fits

```
models %>% select(continent, augmented) %>% unnest(augmented) %>%
  ggplot(aes(x=log(gdpPercap), y=fitted, color=continent, fill=continent)) +
  geom_point(size = 0.8, alpha=0.5) +
  geom_smooth(method = "lm", alpha=0.5) +
  ylab("Fitted Life Expectancy")
```

The slope for the Americas is noticeably larger than for other continents



46

Nice tables in R

- Not a ggplot topic, but it is useful to know that you can also produce beautiful tables in R
- There are many packages for this: See the CRAN Task View on Reproducible Research, <https://cran.r-project.org/web/views/ReproducibleResearch.html>
 - xtable: Exports tables to LaTeX or HTML, with lots of control
 - gt: the ggplot of tables!
 - flextable: similar to gt, but with themes
 - stargazer: Well-formatted model summary tables, side-by-side
 - apaStyle: Generate APA Tables for MS Word

47

Tables in R: xtable

Just a few examples, stolen from xtable: vignette("xtableGallery.pdf")

```
fm1 <- aov(tlimth ~ sex + ethnicity + grade + disadvg, data = tli)
xtable(fm1)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
sex	1	75.37	75.37	0.38	0.5417
ethnicity	3	2572.15	857.38	4.27	0.0072
grade	1	36.31	36.31	0.18	0.6717
disadvg	1	59.30	59.30	0.30	0.5882
Residuals	93	18682.87	200.89		

```
fm3 <- glm(disadvg ~ ethnicity*grade, data = tli, family = binomial)
xtable(fm3)
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	3.1888	1.5966	2.00	0.0458
ethnicityHISPANIC	-0.2848	2.4808	-0.11	0.9086
ethnicityOTHER	212.1701	22122.7093	0.01	0.9923
ethnicityWHITE	-8.8150	3.3355	-2.64	0.0082
grade	-0.5308	0.2892	-1.84	0.0665
ethnicityHISPANIC:grade	0.2448	0.4357	0.56	0.5742
ethnicityOTHER:grade	-32.6014	3393.4687	-0.01	0.9923
ethnicityWHITE:grade	1.0171	0.5185	1.96	0.0498


Too many decimals are used here, but you can control all that

48

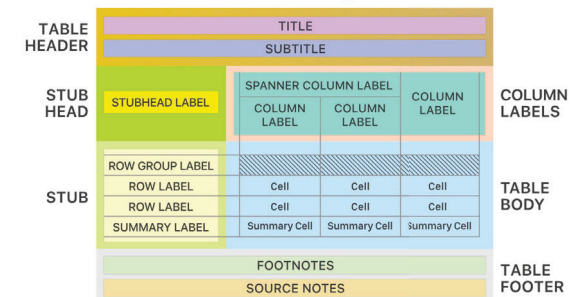
Tables with {gt}

- The {gt} package aims to provide a grammar of tables just as ggplot2 does for graphs
 - Designed to be simple to use, yet powerful

```
iris %>% gt()
```



The Parts of a gt Table




49

{gt} workflow

- Data table -> gt_tbl object

```
iris %>% gt() %>%
  tab_header(...) %>%
  tab_options(...) %>% gtsave()
```

A Typical gt Workflow



50

```
iris_tab <-
iris %>%
  slice(1:5) %>%
  gt() %>%
  tab_header(
    title = "Anderson's Iris Data",
    subtitle = "(Collected in ...)")
```

Sample 5 rows
pipe to gt()

add header

Anderson's Iris Data (Collected in the Gaspé Peninsula)				
Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
6.4	3.1	5.5	1.8	virginica
6.9	3.1	5.4	2.1	virginica
4.4	3.2	1.3	0.2	setosa
7.7	2.8	6.7	2.0	virginica
5.0	2.3	3.3	1.0	versicolor

51

```
iris_tab <-
iris %>%
  slice(1:5) %>%
  gt() %>%
  tab_header(
    title = "Anderson's Iris Data",
    subtitle = "(Collected in ...)")
```

```
iris_tab <-
iris_tab %>%
  tab_spanner(
    label = "Sepal",
    columns = c(Sepal.Length,
                Sepal.Width)) %>%
  tab_spanner(
    label = "Petal",
    columns = c(Petal.Length,
                Petal.Width)) %>%
```

Add table column spanning headers

Anderson's Iris Data (Collected in the Gaspé Peninsula)				
Sepal		Petal		
Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
6.4	3.1	5.5	1.8	virginica
6.9	3.1	5.4	2.1	virginica
4.4	3.2	1.3	0.2	setosa
7.7	2.8	6.7	2.0	virginica
5.0	2.3	3.3	1.0	versicolor

52

```
iris_tab <-
iris %>%
  slice(1:5) %>%
  gt() %>%
  tab_header(
    title = "Anderson's Iris Data",
    subtitle = "(Collected in ...)")
```

```
iris_tab <- iris_tab %>%
  cols_label(
    Sepal.Length = "Length",
    Sepal.Width = "Width",
    Petal.Length = "Length",
    Petal.Width = "Width") %>%
  tab_options(
    heading.background.color = "#c6dbef",
    column_labels.background.color = "#edf8fb")
```

Re-label columns

Colorize headings

```
iris_tab <-
iris_tab %>%
  tab_spanner(
    label = "Sepal",
    columns = c(Sepal.Length,
                Sepal.Width)) %>%
  tab_spanner(
    label = "Petal",
    columns = c(Petal.Length,
                Petal.Width)) %>%
```

Anderson's Iris Data (Collected in the Gaspé Peninsula)				
Sepal		Petal		
Length	Width	Length	Width	Species
6.4	3.1	5.5	1.8	virginica
6.9	3.1	5.4	2.1	virginica
4.4	3.2	1.3	0.2	setosa
7.7	2.8	6.7	2.0	virginica
5.0	2.3	3.3	1.0	versicolor

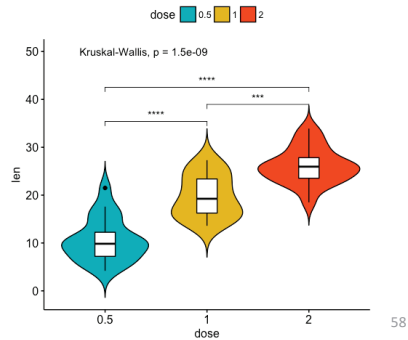
53

ggpubr

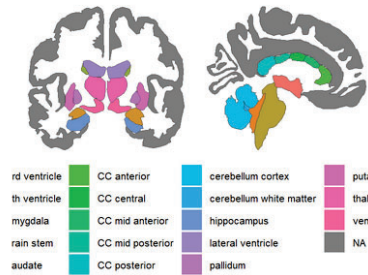
The **ggpubr** package provides some easy-to-use functions for creating and customizing publication ready plots.

```
ggviolin(df, x = "dose", y = "len", fill = "dose",
  palette = c("#00AFBB", "#E7B800", "#FC4E07"),
  add = "boxplot", add.params = list(fill = "white")) +
  stat_compare_means(comparisons = my_comparisons, label = "p.signif") +
  stat_compare_means(label.y = 50)
```

see the examples at
<http://www.sthda.com/english/rpkgs/ggpubr/>



ggseg: plotting brain atlases



```
# install.packages("remotes")
# remotes::install_github("LCBC-UiO/ggseg")
library(ggseg)
# install.packages("ggplot2")
library(ggplot2)

ggplot() +
  geom_brain(atlas = aseg) +
  theme_void() +
  theme(legend.position = "bottom",
    legend.text = element_text(size = 8)) +
  guides(fill = guide_legend(ncol = 4))
```