

ggplot2: Going further in the tidyverse


Michael Friendly

Psych 6135

<https://friendly.github.io/6135/>

A larger view: Data science


- Data science treats statistics & data visualization as parts of a larger process
 - Data import: text files, data bases, web scraping, ...
 - Data cleaning → “tidy data”
 - Model building & visualization
 - Reproducible report writing



2

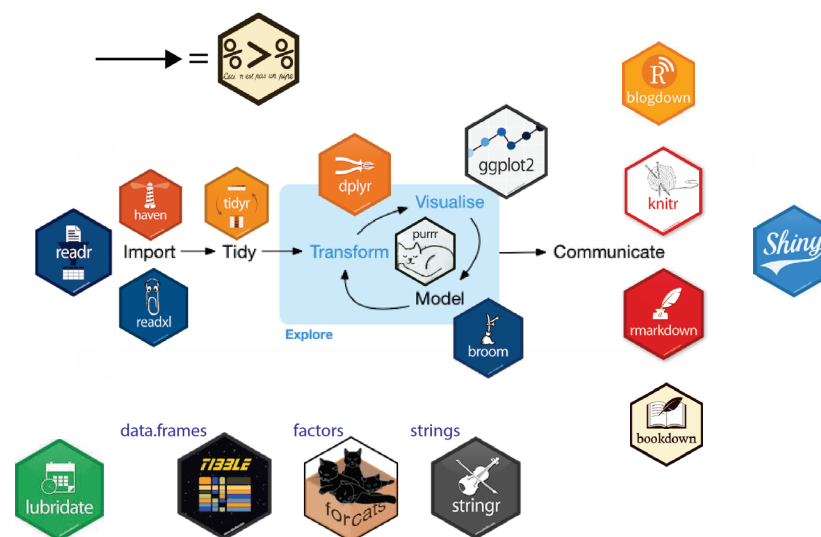


The tidyverse of R packages



3

The tidyverse expands



4

Topics

- Data import / export
- Data wrangling: getting your data into shape
 - dplyr & tidyr
 - pipes: %>%
 - grouping & summarizing
 - Example: NASA data on solar radiation
- Visualizing models: broom
 - Example: gapminder data
- Bootstrapping
- ggplot2 extensions
- tables in R

5

Ready for some heavy lifting?



6

Data Import / Export

- The readr package is the modern, tidy way to import and export data
 - Tabular data:
 - comma delimited (read.csv)
 - any other delimiters (";" = read.csv2; <tab> = read_tsv)
 - Data types:
 - specify column types or let functions guess
- Other data formats

package	Data types
haven	SAS, SPSS, Stata
readxl	Excel files (.xls and .xlsx)
DBI	Databases (SQL, ...)
rvest	HTML (web scraping)

7

Data Import: RStudio

Environment History Connections Build

Import Dataset

From Text (base)...
From Text (readr)...
From Excel...
From SPSS...
From SAS...
From Stata...

Environment:

file:

Import Text Data

File/URL:
C:/Users/friendly/Dropbox/Documents/6135/R/drugs.txt

Data Preview:

subject (character)	drug1 (double)	drug2 (double)	drug3 (double)	drug4 (double)
sub1	20	34	38	44
sub2	16	28	30	34
sub3	14	28	26	30
sub4	18	20	24	30
sub5	10	18	14	22

options:

Import Options:

Name: drugs
Skip: 0

☒ First Row as Names
☒ Trim Spaces
☒ Open Data Viewer

Delimiter: Whitespace
Quotes: Default
Locale: Configure...

Escape: None
Comment: Default
NA: Default

code:




```
library(readr)
drugs <- read_table2("R/drugs.txt")
view(drugs)
```

Reading rectangular data using readr

8

Data transformation tools

Some common data types can be messy when imported. Tidy tools are there to help

dates/times	lubridate	read dates/times in various formats; extract components	
factors	forcats	Change order of levels, drop levels, combine levels	
strings	stringr	detect matches, subset, replace	

9



lubridate: Dates & times

PARSE DATE-TIMES (Convert strings or numbers to date-times)

1. Identify the order of the year (y), month (m), day (d), hour (h), minute (m) and second (s) elements in your data.
2. Use the function below whose name replicates the order. Each accepts a wide variety of input formats.

2017-11-28T14:02:00 ymd_hms(), ymd_hm(), ymd_h().
ymd_hms("2017-11-28T14:02:00")

2017-22-12 10:00:00 ydm_hms(), ydm_hm(), ydm_h().
ydm_hms("2017-22-12 10:00:00")

11/28/2017 1:02:03 mdy_hms(), mdy_hm(), mdy_h().
mdy_hms("11/28/2017 1:02:03")

1 Jan 2017 23:59:59 dmy_hms(), dmy_hm(), dmy_h().
dmy_hms("1 Jan 2017 23:59:59")

20170131 ymd(), ydm(). ymd(20170131)

July 4th, 2000 mdy(), myd(). mdy("July 4th, 2000")

4th of July '99 dmy(), dym(). dmy("4th of July '99")

2001: Q3 yq() Q for quarter. yq("2001: Q3")

2:01 hms:hms() Also lubridate::hms(),
hm() and ms(), which return
periods. hms:hms(sec = 0, min = 1,
hours = 2)

GET AND SET COMPONENTS

Use an accessor function to get a component.
Assign into an accessor function to change a component in place.

d ## "2017-11-28"
day(d) ## 28
day(d) <- 1
d ## "2017-11-01"

2018-01-31 11:59:59 date(x) Date component. date(dt)

2018-01-31 11:59:59 year(x) Year. year(dt)
isoyear(x) The ISO 8601 year.
epiyear(x) Epidemiological year.

2018-01-31 11:59:59 month(x, label, abbr) Month.
month(dt)

2018-01-31 11:59:59 day(x) Day of month. day(dt)
wday(x, label, abbr) Day of week.
qday(x) Day of quarter.

2018-01-31 11:59:59 hour(x) Hour. hour(dt)

2018-01-31 11:59:59 minute(x) Minutes. minute(dt)

2018-01-31 11:59:59 second(x) Seconds. second(dt)

week(x) Week of the year. week(dt)
isoweek() ISO 8601 week.
epiweek() Epidemiological week.

Learn more at: <http://lubridate.tidyverse.org>

10



stringr: Manipulating strings

Detect Matches

str_detect(string, pattern) Detect the presence of a pattern match in a string.
str_detect(fruit, "a")

str_which(string, pattern) Find the indexes of strings that contain a pattern match.
str_which(fruit, "a")

str_count(string, pattern) Count the number of matches in a string.
str_count(fruit, "a")

str_locate(string, pattern) Locate the positions of pattern matches in a string. Also
str_locate_all, str_locate(fruit, "a")

Subset Strings

str_sub(string, start = 1L, end = -1L) Extract substrings from a character vector.
str_sub(fruit, 1, 3); str_sub(fruit, -2)

str_subset(string, pattern) Return only the strings that contain a pattern match.
str_subset(fruit, "b")

str_extract(string, pattern) Return the first pattern match found in each string, as a vector. Also str_extract_all to return every pattern match.
str_extract(fruit, "[aeiou]")

str_match(string, pattern) Return the first pattern match found in each string, as a matrix with a column for each (1) group in pattern. Also str_match_all.
str_match(sentences, "[a-z]+")

Mutate Strings

str_sub(<string>, <value>) Replace substrings by identifying the substrings with str_sub() and assigning into the results.
str_sub(fruit, 1, 3) <- "str"

str_replace(string, pattern, replacement) Replace the first matched pattern in each string. str_replace(fruit, "a", "-")

str_replace_all(string, pattern, replacement) Replace all matched patterns in each string. str_replace_all(fruit, "a", "-")

str_to_lower(string, locale = "en") Convert strings to lower case.
str_to_lower(sentences)

str_to_upper(string, locale = "en") Convert strings to upper case.
str_to_upper(sentences)

Join and Split

str_c(..., sep = "", collapse = NULL) Join multiple strings into a single string.
str_c(letters, LETTERS)

str_c(..., sep = "", collapse = "") Collapse a vector of strings into a single string.
str_c(letters, collapse = "")

str_dup(string, times) Repeat strings times times.
str_dup(fruit, times = 2)

str_split_fixed(string, pattern, n) Split a vector of strings into a matrix of substrings (splitting at occurrences of a pattern match). Also str_split to return a list of substrings.
str_split(fruit, "[aeiou]")

str_glue(..., sep = "", env = parent.frame()) Create a string from strings and (expressions) to evaluate. str_glue("Pi is {pi}")

Learn more at: <http://stringr.tidyverse.org>

11



forcats: Working with factors

R represents categorical variables as factors, useful for analysis (e.g., ANOVA)
In graphics, we often want to recode levels or reorder them

Factors

R represents categorical data with factors. A factor is an integer vector with a levels attribute that stores a set of mappings between integers and categorical values. When you view a factor, R displays not the integers, but the values associated with them.

Create a factor with factor()
factor(x = character(), levels = levels, exclude = NA, ordered = is.ordered(x), nmax = NA) Convert a vector to a factor. Also as.factor.
f <- factor(c("a", "b", "b", "a", "a"), levels = c("a", "b", "c"))

Return its levels with levels()
levels(x) Return/set the levels of a factor. levels(f); levels(f) <- c("x", "y", "z")

Use unclass() to see its structure

Inspect Factors

fact_count(f, sort = FALSE)
Count the number of values with each level. fact_count(f)

Change the order of levels

fact_relevel(f, ..., after = 0L)
Manually reorder factor levels.
fact_relevel(f, c("b", "a", "c"))

fact_freq(f, ordered = NA)
Reorder levels by the frequency in which they appear in the data (highest frequency first).
f <- factor(c("a", "c", "c", "a", "a"))
fact_freq(f)

fact_inorder(f, ordered = NA)
Reorder levels by order in which they appear in the data.
fact_inorder(f)

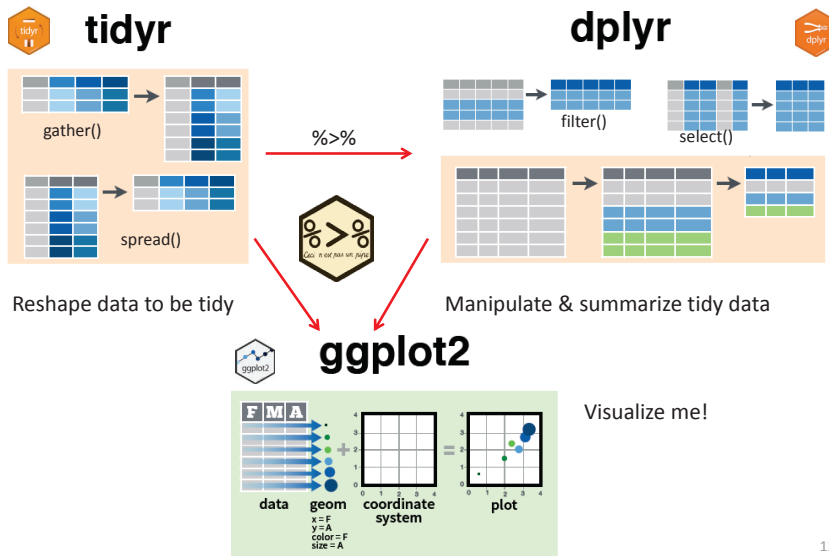
fact_rev(f) Reverse level order.
f <- factor(c("a", "b", "b", "a", "a"))
fact_rev(f)

fact_shift(f) Shift levels to left or right, wrapping around end.
fact_shift(f)

Learn more at: <http://forcats.tidyverse.org>

12

Tidy tools: overview



13


Tidy operations

Reshape long to wide

synonym: **tidyr::pivot_longer()**

Reshape long to wide

synonym: **tidyr::pivot_longer()**



Separate parts of a value into several variables

Join related variables into one


14

Data wrangling with dplyr & tidyr

What is Tidy Data?

A dataset is said to be tidy if:

- observations are in **rows**
- variables are in **columns**
- each value is in its own **cell**.



A "messy" dataset: Survey of income by religion from Pew Research

- Values of **income** are in separate columns, not one variable
- Column headers are **values**, not variable names
- Cell values are frequencies--- **implicit**, not explicit

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	\$50-75k
Agnostic	27	34	60	81	76	137
Atheist	12	27	37	52	35	70
Buddhist	27	21	30	34	33	58
Catholic	418	617	732	670	638	1116

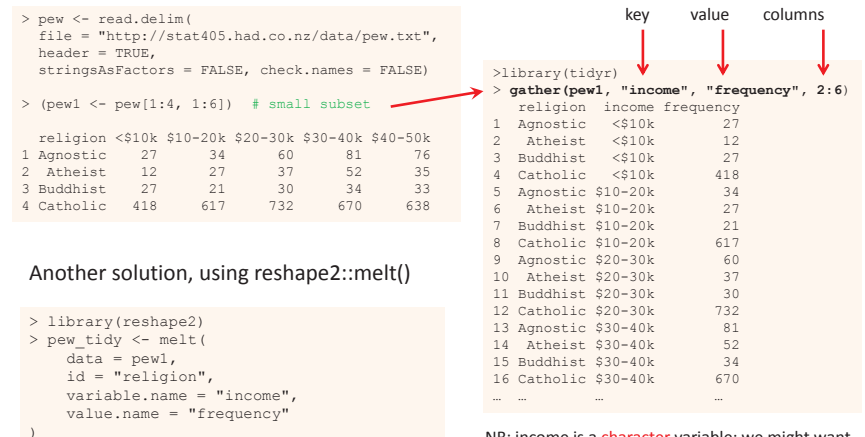
This organization is easy in Excel

But, this makes data analysis and graphing hard

15

Tidying: reshaping wide to long

We can tidy the data by reshaping from wide to long format using **tidyr::gather()**



16



Using pipes: %>%

- R is a functional language

- This means that $f(x)$ returns a value, as in $y \leftarrow f(x)$
- That value can be passed to another function: $g(f(x))$
- And so on: $h(g(f(x)))$

```
> x <- c(0.109, 0.359, 0.63, 0.996, 0.515, 0.142)
> exp(diff(log(x)))
[1] 3.29 1.75 1.58 0.52 0.28
```

- This gets messy and hard to read, unless you break it down step by step

```
> # Compute the logarithm of `x`, calculate lagged differences,
> # return the exponential function of the result
> log(x)
[1] -2.216 -1.024 -0.462 -0.004 -0.664 -1.952
> diff(log(x)) #calculate lagged diffs
[1] 1.19 0.56 0.46 -0.66 -1.29
> exp(diff(log(x))) # convert back to original scale
[1] 3.29 1.75 1.58 0.52 0.28
```

17

Using pipes: %>%

- Pipes (%>%) change the syntax to make this easier

```
> # use pipes
> x %>% log() %>% diff() %>% exp()
[1] 3.29 1.75 1.58 0.52 0.28
```

- Basic rules

- $x \%>\% f()$ passes object on left hand side as first argument (or . argument) of function on right hand side
 - $x \%>\% f()$ is the same as $f(x)$
 - $x \%>\% f(y)$ is the same as $f(x, y)$
 - $y \%>\% f(x, ., z)$ is the same as $f(x, y, z)$
- $x \%<>\% f()$ does the same, but assigns the result to x
 - Shortcut for $x \leftarrow x \%>\% f()$

18



Using pipes: %>% ggplot()




For the Pew data, mutate income into an ordered factor and make a ggplot

```
pew1 %>%
  gather("income", "frequency", 2:6) %>% # reshape
  mutate(income = ordered(income, levels=unique(income))) %>% # make ordered
  ggplot(aes(x=income, fill=religion)) + # plot
  geom_bar(aes(weight=frequency)) # as freq bars
```

`mutate()` calculates or transforms column variables
ordered(income) levels are now ordered appropriately.

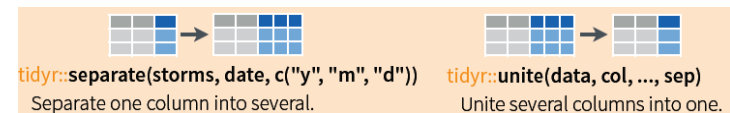
The result is piped to `ggplot()`



19

Tidying: separate() and unite()

It sometimes happens that several variables are crammed into one column, or parts of one variable are split across multiple columns



For example, for the pew data, we might want separate income into low & high

```
pew_long %>%
  mutate(inc = gsub("[\\$k]", "", income)) %>%
  mutate(inc = gsub("<", "0-", inc)) %>%
  separate(inc, c("low", "high", "-")) %>%
  head()
```

	religion	income	frequency	low	high
1	Agnostic	<\$10k	27	0	10
2	Atheist	<\$10k	12	0	10
3	Buddhist	<\$10k	27	0	10
4	Catholic	<\$10k	418	0	10
5	Agnostic	\$10-20k	34	10	20
6	Atheist	\$10-20k	27	10	20

20

dplyr: Subset observations (rows)

dplyr implements a variety of verbs to select a subset of observations from a dataset



In a pipe expression, omit the dataset name

dplyr::filter(iris, Sepal.Length > 7)
Extract rows that meet logical criteria.

dplyr::distinct(iris)
Remove duplicate rows.

dplyr::sample_frac(iris, 0.5, replace = TRUE)
Randomly select fraction of rows.

dplyr::sample_n(iris, 10, replace = TRUE)
Randomly select n rows.

dplyr::slice(iris, 10:15)
Select rows by position.

dplyr::top_n(storms, 2, date)
Select and order top n entries (by group if grouped data).

```
iris %>% filter(Sepal.Length > 7)
iris %>% filter(Species == "setosa")

iris %>% sample_n(10)
iris %>% slice(1:50) # setosa
```

21

dplyr: Subset variables (columns)



dplyr::select(iris, Sepal.Width, Petal.Length, Species)

Select columns by name or helper function.

Many helper functions in dplyr allow selection by a **function** of variable names:

select(iris, contains(" "))
Select columns whose name contains a character string.

select(iris, ends_with("Length"))
Select columns whose name ends with a character string.

select(iris, everything())
Select every column.

select(iris, matches("t."))
Select columns whose name matches a regular expression.

select(iris, num_range("x", 1:5))
Select columns named x1, x2, x3, x4, x5.

select(iris, one_of(c("Species", "Genus")))
Select columns whose names are in a group of names.

select(iris, starts_with("Sepal"))
Select columns whose name starts with a character string.

select(iris, Sepal.Length:Petal.Width)
Select all columns between Sepal.Length and Petal.Width (inclusive).

select(iris, -Species)
Select all columns except Species.


22

dplyr: group_by() and summarise()

Fundamental operations in data munging are:

- **grouping** a dataset by one or more variables
- calculating one or more **summary** measures
- **ungrouping**: expand to an ungrouped copy, if needed

data group summarise



```
mtcars %>%
  group_by(cyl) %>%
  summarise(avg = mean(mpg))
```

data avg



```
mtcars %>%
  group_by(cyl) %>%
  summarise(avg = mean(mpg)) %>%
  ungroup()
```

23

Example: NASA data on solar radiation



Surface meteorology and Solar Energy

A renewable energy resource web site (release 6.0)
sponsored by NASA's Applied Science Program in the Science Mission Directorate
developed by POWER: Prediction of Worldwide Energy Resource Project

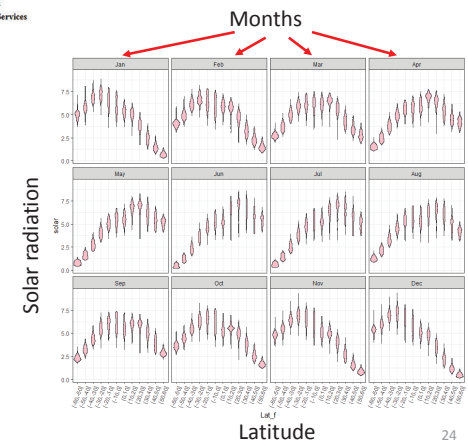


- over 200 satellite-derived meteorology and solar energy parameters
- monthly averaged from 22 years of data
- data tables for a particular location
- GIS Web Mapping Application & Services

How does solar radiation vary with latitude, over months of the year?

How to make this plot?

Q:
what are the basic plot elements?




24

NASA data: solar radiation

This is easy to do for the total **Annual** solar radiation, a column in the data

```
> str(nasa)
'data.frame': 64800 obs. of 15 variables:
 $ Lat: int -90 -90 -90 -90 -90 -90 -90 -90 ...
 $ Lon: int -180 -179 -178 -177 -176 -175 -174 -173 -172 -171 ...
 $ Jan: num 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 ...
 $ Feb: num 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 ...
 $ Mar: num 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 ...
 $ Apr: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ May: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Jun: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Jul: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Aug: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Sep: num 0.10 0.10 0.10 0.10 0.10 0.10 0.10 0.10 ...
 $ Oct: num 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 ...
 $ Nov: num 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 ...
 $ Dec: num 11.11 11.11 11.11 11.11 11.11 11.11 11.11 11.11 ...
 $ Ann: num 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 ...
```



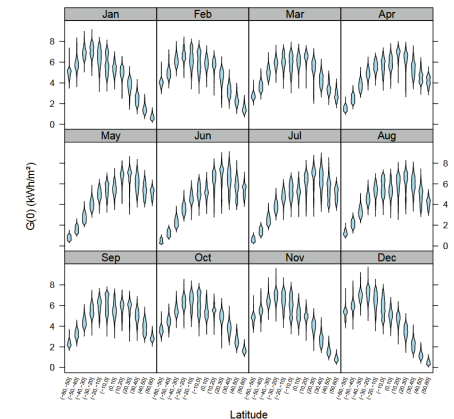
```
nasa %>%
  filter(abs(Lat) < 60) %>%
  mutate(Latf = cut(Lat, pretty(Lat, n=10))) %>%
  ggplot(aes(x=Latf, y=Ann)) +
    geom_violin(fill="pink", alpha=0.3) +
    labs(x="Latitude", y="Solar radiation G(0) (kWh/m²)")
```

25

Faceting & tidy data

This is complicated to do for the separate months, because the data structure is **untidy**--- months were in separate variables (wide format)

```
> str(nasa)
'data.frame': 64800 obs. of 15 variables:
 $ Lat: int -90 -90 -90 -90 -90 -90 -90 -90 ...
 $ Lon: int -180 -179 -178 -177 -176 -175 -174 -173 -172 -171 ...
 $ Jan: num 9.63 9.63 9.63 9.63 9.63 9.63 9.63 9.63 ...
 $ Feb: num 5.28 5.28 5.28 5.28 5.28 5.28 5.28 5.28 ...
 $ Mar: num 0.75 0.75 0.75 0.75 0.75 0.75 0.75 0.75 ...
 $ Apr: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ May: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Jun: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Jul: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Aug: num 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 ...
 $ Sep: num 0.10 0.10 0.10 0.10 0.10 0.10 0.10 0.10 ...
 $ Oct: num 3.24 3.24 3.24 3.24 3.24 3.24 3.24 3.24 ...
 $ Nov: num 8.28 8.28 8.28 8.28 8.28 8.28 8.28 8.28 ...
 $ Dec: num 11.11 11.11 11.11 11.11 11.11 11.11 11.11 11.11 ...
 $ Ann: num 3.19 3.19 3.19 3.19 3.19 3.19 3.19 3.19 ...
```



26

tidying the data

To plot solar radiation against latitude by month (separate panels), we need to:

- remove the Ann column
- reshape the data to long format, so solar is all in one column

```
library(tidyr)
library(dplyr)
library(ggplot2)
```

```
nasa_long <- nasa %>%
  select(-Ann) %>%
  gather(month, solar, Jan:Dec, factor_key=TRUE) %>%
  filter( abs(Lat) < 60 ) %>%
  mutate( Lat_f = cut(Lat, pretty(Lat, 12)))
```

%>% "pipes" data to the next stage

select() extracts or drops columns
gather() collapses columns into key-value pairs
filter() subsets observations
mutate() creates new variables

27

tidying the data

```
> str(nasa_long)
'data.frame': 514080 obs. of 5 variables:
 $ Lat : int -59 -59 -59 -59 -59 -59 -59 -59 ...
 $ Lon : int -180 -179 -178 -177 -176 -175 -174 -173 -172 -171 ...
 $ month: Factor w/ 12 levels "Jan","Feb","Mar",...: 1 1 1 1 1 1 1 1 ...
 $ solar: num 5.19 5.19 5.25 5.25 5.17 5.17 5.15 5.15 5.15 5.15 ...
 $ Lat_f: Factor w/ 12 levels "(-60,-50]",...: 1 1 1 1 1 1 1 1 ...
```

For ease of plotting, I created a factor version of Lat with 12 levels

```
> head(nasa_long)
  Lat Lon month solar Lat_f
1 -59 -180 Jan 5.19 (-60,-50]
2 -59 -179 Jan 5.19 (-60,-50]
3 -59 -178 Jan 5.25 (-60,-50]
4 -59 -177 Jan 5.25 (-60,-50]
5 -59 -176 Jan 5.17 (-60,-50]
6 -59 -175 Jan 5.17 (-60,-50]
```

The data are now in a form where I can plot solar against Lat or Lat_f and facet by month

28


plotting the tidy data

Using `geom_violin()` shows the shapes of the distributions for levels of `Lat_f`

```
ggplot(nasa_long, aes(x=Lat_f, y=solar)) +
  geom_violin(fill="pink") +
  facet_wrap(~ month) +
  theme_bw() +
  theme(axis.text.x =
    element_text(angle = 70,
      hjust = 1))
```

`facet_wrap(~month)` does the right thing

I had to adjust the x-axis labels for `Lat_f` to avoid overplotting



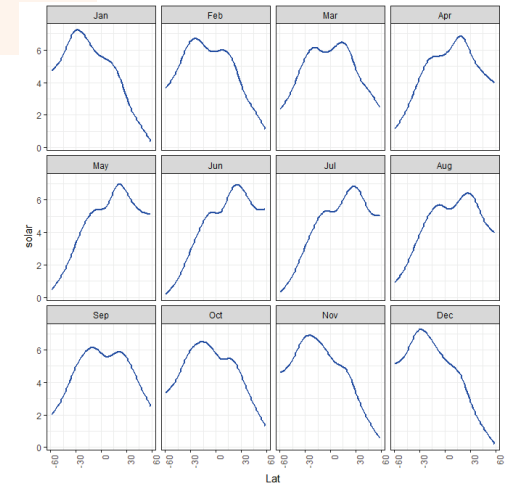
plotting the tidy data: smoothing

```
ggplot(nasa_long, aes(x=Lat, y=solar)) +
  geom_smooth(color="blue") +
  facet_wrap(~ month) +
  theme_bw()
```

Here we treat `Lat` as quantitative. `geom_smooth()` uses method = "gam" here because of large n

The variation in the smoothed trends over the year suggest quite lawful behavior

Can we express this as a statistical model ?



build a model

What we saw in the plot suggests a **generalized additive model**, with a smooth, $s(\text{Lat})$

```
library(mgcv)
nasa.gam <- gam(solar ~ Lon + month + s(Lat), data=nasa_long)
summary(nasa.gam)
```

Family: gaussian
Link function: identity

Formula:
 $\text{solar} \sim \text{Lon} + \text{month} + s(\text{Lat})$

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.691e+00	6.833e-03	686.409	< 2e-16 ***
Lon	-1.713e-04	1.898e-05	-9.022	< 2e-16 ***
monthFeb	1.195e-01	9.664e-03	12.364	< 2e-16 ***
...
monthDec	-8.046e-02	9.664e-03	-8.326	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The violin plots suggest that variance is not constant. I'm ignoring this here by using the default gaussian model.

Model terms:

- `Lon` wasn't included before
- `month` is a factor, for the plots
- $s(\text{Lat})$ fits a smoothed term in latitude, averaged over other factors


There are other model choices, but it is useful to visualize what we have done so far

visualize the model


Effect plots show the fitted relationship between the response and model terms, averaged over other predictors.

The `mgcv` package has its own versions of these.


```
plot(nasa.gam, cex.lab=1.25)
termplot(nasa.gam, terms="month", se=TRUE, lwd.term=3, lwd.se=2, cex.lab=1.25)
termplot(nasa.gam, terms="Lon", se=TRUE, lwd.term=3, lwd.se=2, cex.lab=1.25)
```



why the dip at the equator?



effect of longitude is very small, but maybe interpretable



month should be modeled as a time variable

Visualizing models

- R modeling functions [`lm()`, `glm()`, ...] return model objects, but these are “messy”
 - extracting coefficients takes several steps: `data.frame(coef(mymod))`
 - some info (R^2 , F , p -value) is computed in `print()` method, not stored
 - can’t easily combine models
- Some have associated plotting functions
 - `plot(model)`: diagnostic plots
 - `car` package: many model plot methods
 - `effects` package: plot effects for model terms
- But what if you want to:
 - make a table of model summary statistics
 - fit a **collection** of models, compare, summarize or visualize them?

33



broom: visualizing models

- The broom package turns model objects into tidy data frames
 - glance**(models) extracts model-level summary statistics (R^2 , df, AIC, BIC)
 - tidy**(models) extracts coefficients, SE, p-values
 - augment**(models) extracts observation-level info (residuals, ...)

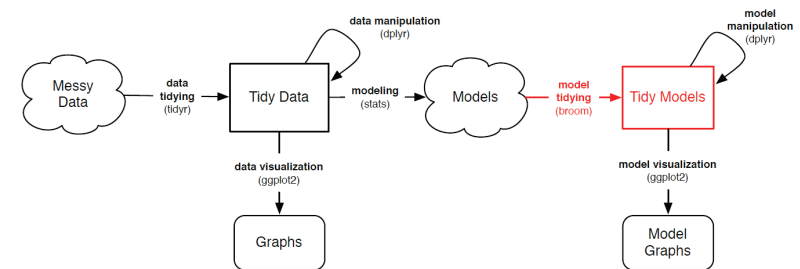


Image from: https://opr.princeton.edu/workshops/Downloads/2016Jan_BroomRobinson.pdf

34


Example: gapminder data

```
ggplot(aes(x = log(gdpPercap), y=lifeExp, color=continent), data=gapminder) +
  geom_point() +
  geom_smooth(method = "loess")
```

How to model this?

How to extract & plot model statistics?

How to fit & display multiple models for subsets?



35

Example: gapminder data

Predict life expectancy from year, population, GDP and continent:

```
gapmod <- lm(lifeExp ~ year + pop + log(gdpPercap) + continent, data=gapminder)
summary(gapmod)
```

```
Call:
lm(formula = lifeExp ~ year + pop + log(gdpPercap) + continent, data = gapminder)

Residuals:
    Min       1Q   Median       3Q      Max
-24.928  -3.285   0.314   3.699  15.221

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -4.58e+02  1.67e+01  -27.43 < 2e-16 ***
year         2.38e-01  8.61e-03   27.58 < 2e-16 ***
pop          5.40e-09  1.38e-09    3.91 9.5e-05 ***
log(gdpPercap) 5.10e+00  1.60e-01   31.88 < 2e-16 ***
continentAmericas 8.74e+00  4.63e-01   18.86 < 2e-16 ***
continentAsia    6.64e+00  4.09e-01   16.22 < 2e-16 ***
continentEurope  1.23e+01  5.10e-01   24.11 < 2e-16 ***
continentOceania 1.26e+01  1.27e+00    9.88 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.79 on 1696 degrees of freedom
Multiple R-squared:  0.8,    Adjusted R-squared:  0.799
F-statistic: 969 on 7 and 1696 DF,  p-value: <2e-16
```

observation level

component level
(coefficients)

model level

36

glance() gives the **model level** summary statistics

```
> glance(gapmod)
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value df logLik AIC BIC deviance df.residual
1      0.8      0.7992 5.789      969      0 8 -5406 10830 10879      56835      1696
```

tidy() gives the **model component** (term) statistics

```
> tidy(gapmod)
# A tibble: 8 x 5
  term      estimate std.error statistic    p.value
1 (Intercept) -4.585e+02 1.671e+01 -27.433 1.982e-137
2 year        2.376e-01 8.613e-03  27.584 1.122e-138
3 pop         5.403e-09 1.381e-09   3.912 9.496e-05
4 log(gdpPercap) 5.103e+00 1.601e-01  31.876 4.096e-175
5 continentAmericas 8.739e+00 4.635e-01  18.856 3.758e-72
6 continentAsia    6.635e+00 4.091e-01  16.219 4.167e-55
7 continentEurope  1.230e+01 5.102e-01  24.113 1.044e-110
8 continentOceania 1.256e+01 1.270e+00   9.884 1.943e-22
```


augment() gives the **observation level** statistics

```
> augment(gapmod) %>% slice(1:5)
# A tibble: 5 x 12
  lifeExp year pop log.gdpPercap. continent .fitted .se.fit .resid .hat .sigma
<dbl> <int> <int> <dbl> <fct> <dbl> <dbl> <dbl> <dbl> <dbl>
1 28.8 1952 8425333 6.66 Asia 46.0 0.408 -17.1 0.00496 5.78
2 30.3 1957 9240934 6.71 Asia 47.4 0.390 -17.1 0.00454 5.78
3 32.0 1962 10267083 6.75 Asia 48.8 0.376 -16.8 0.00423 5.78
4 34.0 1967 11537966 6.73 Asia 49.9 0.372 -15.9 0.00413 5.78
5 36.1 1972 13079460 6.61 Asia 50.5 0.382 -14.4 0.00435 5.78
# ... with 2 more variables: .cooksd <dbl>, .std.resid <dbl>
```

37

tidyr:: "nest – map – unnest" trick

- In many cases, we want to perform analysis for each subset of a dataset defined by one or more variables
- `dplyr::group_by()`, `summarise()`, `ungroup()` is one way
- `tidyr::nest()`, `purrr::map()`, `tidyr::unnest()` is more general

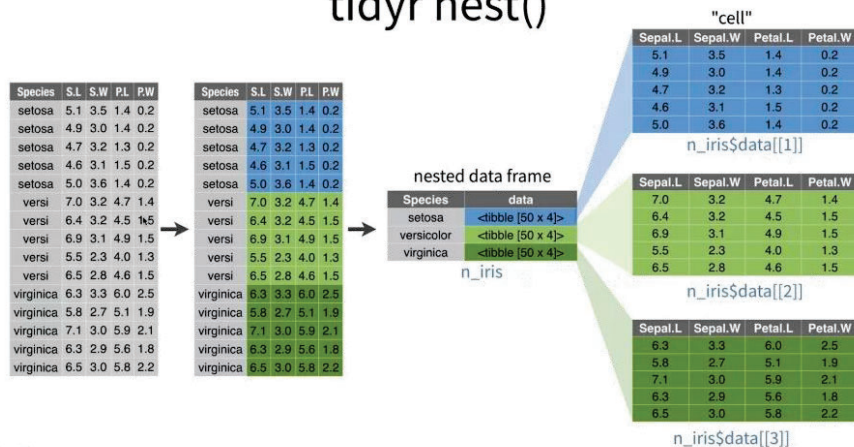


See: https://cran.r-project.org/web/packages/broom/vignettes/broom_and_dplyr.html

38

```
n_iris <- iris %>% group_by(Species) %>% nest() # group by Species, then nest
n_iris <- iris %>% nest(-Species) # nest all other cols
```

tidyr nest()



CC by RStudio

39

tidyr: fitting multiple models

- There may be different effects by continent (GDP x continent interaction)
- What if want to fit (and visualize) a separate model for each continent?
 - nest by continent, then {fit, tidy, glance, augment}

```
models <- gapminder %>%
  filter(continent != "Oceania") %>% # only two countries
  nest(data = -continent) %>%
  mutate(
    fit = map(data, ~ lm(lifeExp ~ year + pop + log(gdpPercap), data = .x)),
    tidied = map(fit, tidy),
    glanced = map(fit, glance),
    augmented = map(fit, augment)
  )
```

What's in this object?

```
names(models)
[1] "continent" "data"      "fit"       "tidied"    "glanced"   "augmented"
```

42

```
# view model summaries
```

```
models %>%
  select(continent, glanced) %>%
  unnest(glanced)
```

Model summary
statistics

```
# A tibble: 4 x 13
  continent r.squared adj.r.squared sigma statistic p.value df logLik AIC
  <fct>      <dbl>      <dbl> <dbl>      <dbl> <dbl> <dbl> <dbl> <dbl>
1 Asia      0.696      0.694 6.56      299. 5.27e-101 3 -1305. 2620.
2 Europe    0.797      0.795 2.46      466. 7.42e-123 3 -833. 1675.
3 Africa    0.500      0.498 6.48      207. 5.90e- 93 3 -2050. 4110.
4 Americas  0.720      0.718 4.97      254. 1.39e- 81 3 -904. 1819.
# ... with 4 more variables: BIC <dbl>, deviance <dbl>, df.residual <int>,
```

```
# model coefficients & tests
```

```
models %>%
  select(continent, tidied) %>%
  unnest(tidied)
```

Coefficients

```
# A tibble: 16 x 6
  continent term estimate std.error statistic p.value
  <fct>      <chr>      <dbl>      <dbl>      <dbl>      <dbl>
1 Asia      (Intercept) -6.20e+2 4.00e+1 -15.5 1.34e-42
2 Asia      year      3.23e-1 2.06e-2 15.7 2.41e-43
3 Asia      pop       5.13e-9 1.66e-9 3.09 2.15e- 3
4 Asia      log(gdpPercap) 5.04e+0 2.76e-1 18.3 2.25e-54
5 Europe    (Intercept) -1.72e+2 1.72e+1 -10.0 4.51e-21
# ...
```

43

```
# observation-level statistics
```

```
models %>%
  select(continent, augmented) %>%
  unnest(augmented)
```

```
# A tibble: 1,680 x 10
  continent lifeExp year pop `log(gdpPercap)` .fitted .hat .sigma .cooks
  <fct>      <dbl> <int> <int> <dbl>      <dbl> <dbl> <dbl> <dbl>
1 Asia      28.8 1952 8425333 6.66      43.7 0.0101 6.53 0.0133
2 Asia      30.3 1957 9240934 6.71      45.6 0.00822 6.53 0.0113
3 Asia      32.0 1962 10267083 6.75      47.4 0.00685 6.53 0.00957
4 Asia      34.0 1967 11537966 6.73      48.9 0.00616 6.53 0.00805
5 Asia      36.1 1972 13079460 6.61      49.9 0.00645 6.54 0.00727
6 Asia      38.4 1977 14880372 6.67      51.9 0.00640 6.54 0.00678
7 Asia      39.9 1982 12881816 6.89      54.6 0.00607 6.53 0.00771
8 Asia      40.8 1987 13867957 6.75      55.5 0.00795 6.53 0.0101
9 Asia      41.7 1992 16317921 6.48      55.8 0.0114 6.53 0.0134
10 Asia     41.8 1997 22227415 6.45      57.3 0.0138 6.53 0.0198
# ... with 1,670 more rows, and 1 more variable: .std.resid <dbl>
```


y predictors \hat{y} diagnostics

44

Visualizing multiple models

One visual summary might be a plot of R^2 values, ordered by continent

```
models %>%
  select(continent, glanced) %>% unnest(glanced) %>%
  ggplot(aes(r.squared, reorder(continent, r.squared))) +
  geom_point(size=4) +
  geom_segment(aes(xend = 0, yend = ..y..)) +
  ylab("Continent")
```



45

Visualizing coefficients


Coefficient plots are often useful, but these are on different scales.

```
models %>% select(continent, tidied) %>% unnest(tidied)
filter(term != "(Intercept)") %>%
mutate(term=factor(term, levels=c("log(gdpPercap)", "year", "pop"))) %>%
ggplot(aes(x=term, y=statistic, color=continent, group=continent)) +
  geom_point(size=5, alpha=0.5) +
  geom_line(size=1.5) +
  geom_hline(yintercept=c(-2, 0, 2), color = c("red", "black", "red")) +
  ylab("t statistic") +
  theme_minimal() + theme(legend.position=c(0.9, 0.8))
```

get model stats
ignore the intercept
reorder terms sensibly
hlines for non-significance

Here, I plot the t -statistics,
 $t = b_{ij}/se(b_{ij})$ for all terms in
all models.

Any values outside $\sim \pm 2$ are
significant, $p < 0.5$!

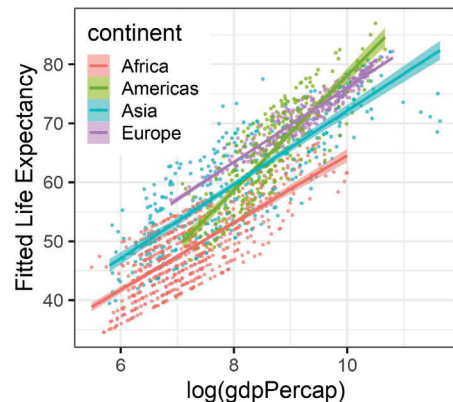


46

Visualizing model fits

```
models %>% select(continent, augmented) %>% unnest(augmented) %>%
  ggplot(aes(x=log(gdpPercap), y=fitted, color=continent, fill=continent)) +
  geom_point(size = 0.8, alpha=0.5) +
  geom_smooth(method = "lm", alpha=0.5) +
  ylab("Fitted Life Expectancy")
```

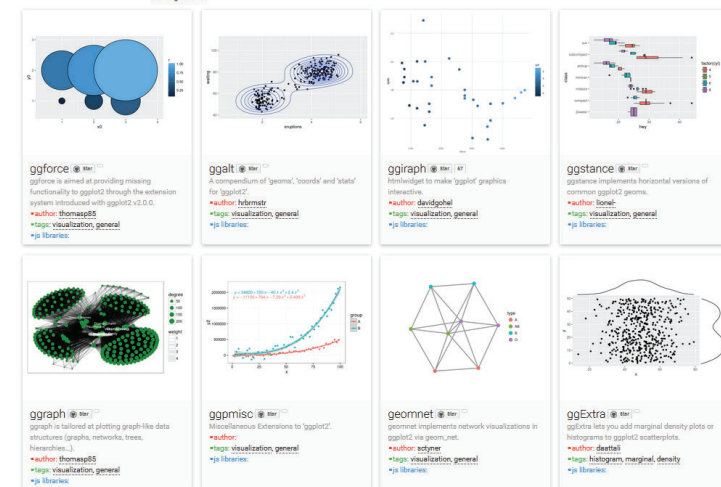
The slope for the Americas is noticeably larger than for other continents



47

ggplot extensions

There are a large number of ggplot extensions. See: <http://www.ggplot2-exts.org/>



48

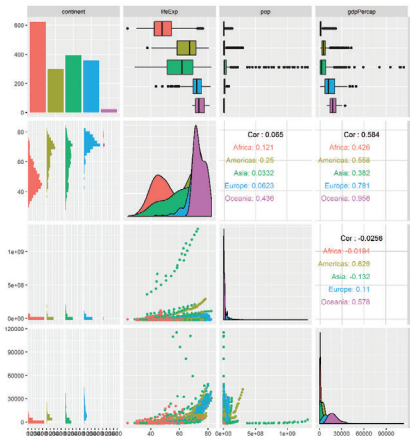
ggplot extensions: GGally

GGally contains a large number of functions that extend ggplot2 to multivariate data

ggpairs() produces generalized scatterplot matrices, with lots of options

```
library(GGally)
library(dplyr)
library(ggplot2)
library(gapminder)

gapminder %>%
  select(-country, -year) %>%
  ggpairs(aes(color=continent))
```

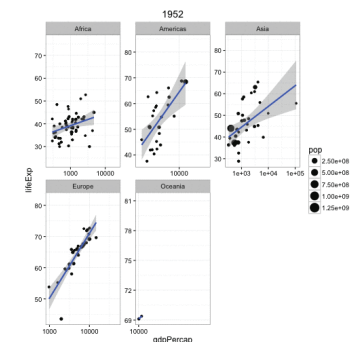


49

ggplot extensions: gganimate

gganimate is a wrapper for the animation package with ggplot2.

It adds a frame= aesthetic, and animates the image as the frame variable changes



```
p5 <- ggplot(gapminder, aes(gdpPercap, lifeExp, size = pop, frame = year)) +
  geom_point() +
  geom_smooth(aes(group = year), method = "lm", show.legend = FALSE) +
  facet_wrap(~continent, scales = "free") +
  scale_x_log10()
```

gganimate(p5)

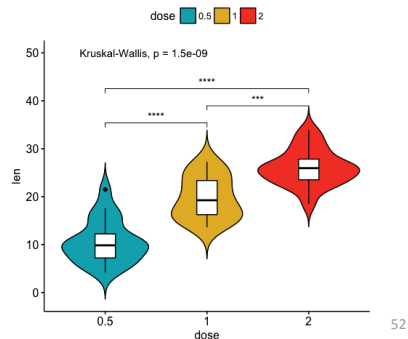
51

ggpubr

The **ggpubr** package provides some easy-to-use functions for creating and customizing publication ready plots.

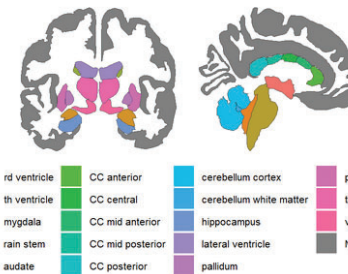
```
ggviolin(df, x = "dose", y = "len", fill = "dose",
  palette = c("#00AFBB", "#E7B800", "#FC4E07"),
  add = "boxplot", add.params = list(fill = "white")) +
  stat_compare_means(comparisons = my_comparisons, label = "p.signif") +
  stat_compare_means(label.y = 50)
```

see the examples at
<http://www.sthda.com/english/rpkgs/ggpubr/>



52

ggseg: plotting brain atlases



```
# install.packages("remotes")
# remotes::install_github("LCBC-UiO/ggseg")
library(ggseg)
# install.packages("ggplot2")
library(ggplot2)

ggplot() +
  geom_brain(atlas = aseg) +
  theme_void() +
  theme(legend.position = "bottom",
    legend.text = element_text(size = 8)) +
  guides(fill = guide_legend(ncol = 4))
```

53

Tables in R

- Not a ggplot topic, but it is useful to know that you can also produce beautiful tables in R
- There are many packages for this: See the CRAN Task View on Reproducible Research, <https://cran.r-project.org/web/views/ReproducibleResearch.html>
 - xtable: Exports tables to LaTeX or HTML, with lots of control
 - stargazer: Well-formatted model summary tables, side-by-side
 - apaStyle: Generate APA Tables for MS Word

55

Tables in R: xtable

Just a few examples, stolen from xtable: vignette("xtableGallery.pdf")

```
fm1 <- aov(tlimth ~ sex + ethnicity + grade + disadv, data = tli)
xtable(fm1)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
sex	1	75.37	75.37	0.38	0.5417
ethnicity	3	2572.15	857.38	4.27	0.0072
grade	1	36.31	36.31	0.18	0.6717
disadv	1	59.30	59.30	0.30	0.5882
Residuals	93	18682.87	200.89		

```
fm3 <- glm(disadv ~ ethnicity*grade, data = tli, family = binomial)
xtable(fm3)
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	3.1888	1.5966	2.00	0.0458
ethnicityHISPANIC	-0.2848	2.4808	-0.11	0.9086
ethnicityOTHER	212.1701	22122.7093	0.01	0.9923
ethnicityWHITE	-8.8150	3.3355	-2.64	0.0082
grade	-0.5308	0.2892	-1.84	0.0665
ethnicityHISPANIC:grade	0.2148	0.4357	0.56	0.5742
ethnicityOTHER:grade	-32.6014	3393.4687	-0.01	0.9923
ethnicityWHITE:grade	1.0171	0.5185	1.96	0.0498

Too many decimals are used here, but you can control all that

56